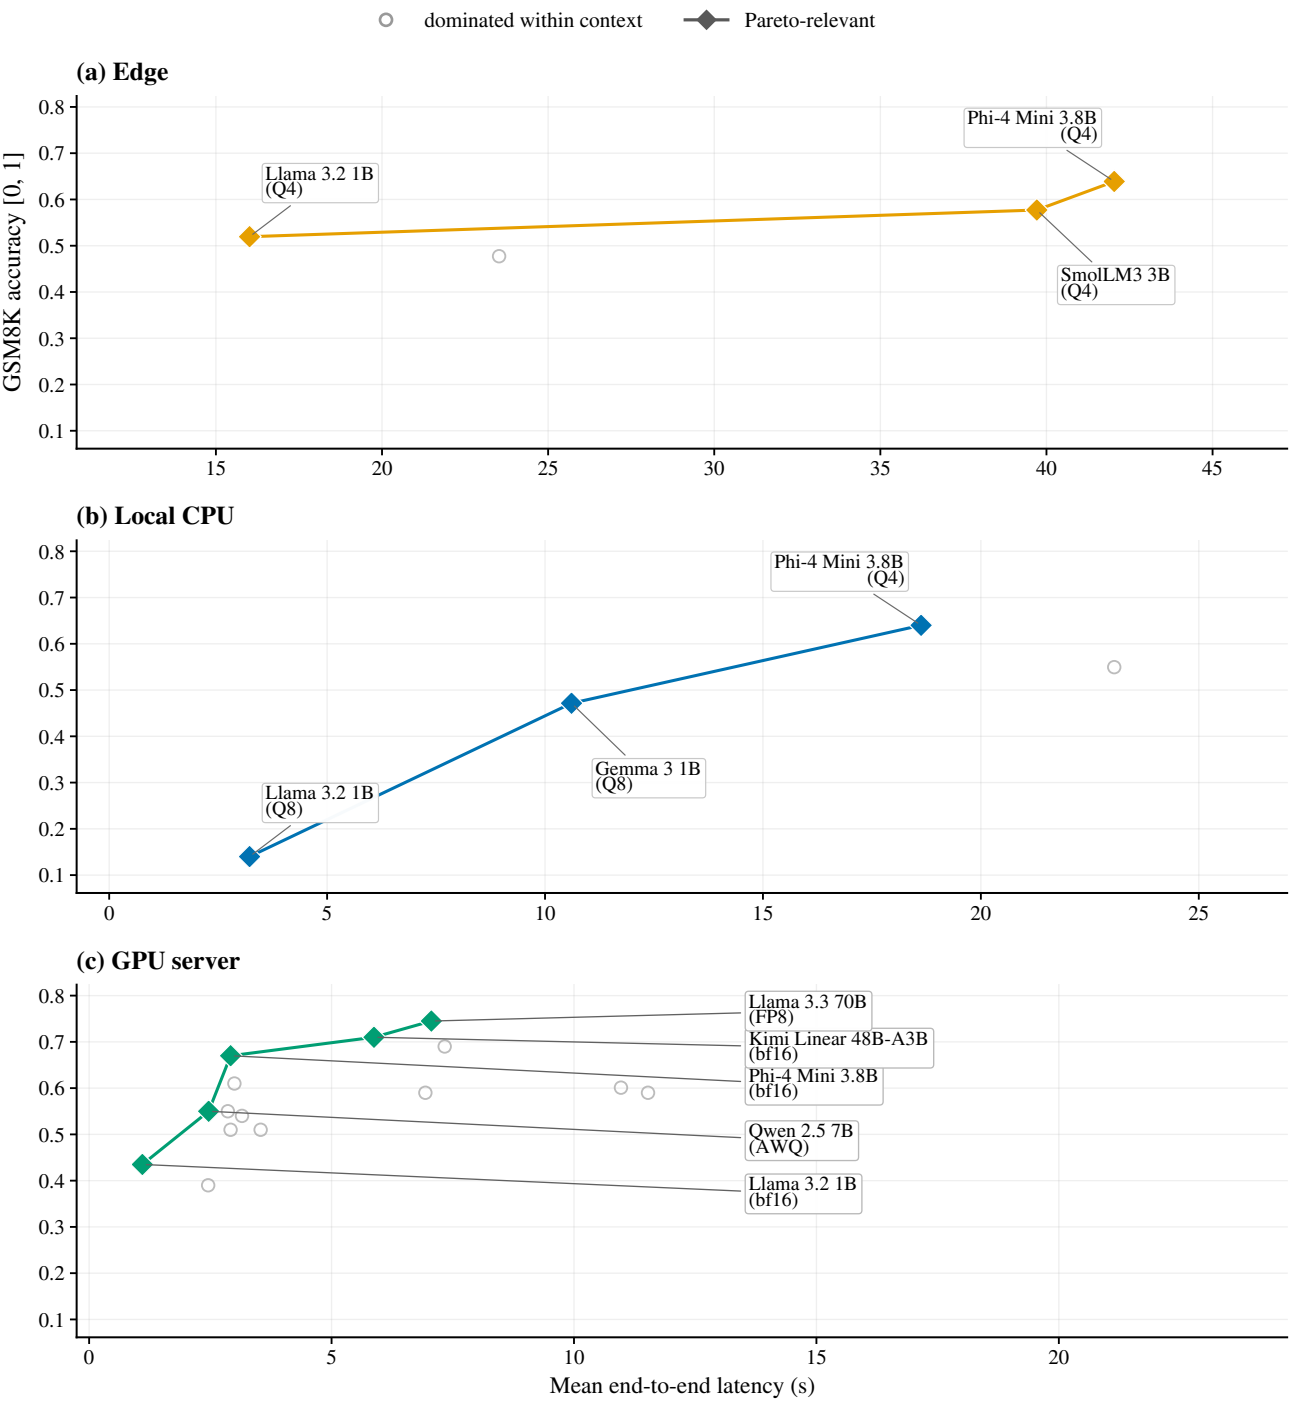
dominated within context
Pareto-relevant
(a) Edge
GSM8K accuracy [0, 1]
Llama 3.2 1B (Q4)
Phi-4 Mini 3.8B (Q4)
SmolLM3 3B (Q4)
(b) Local CPU
Llama 3.2 1B (Q8)
Gemma 3 1B (Q8)
Phi-4 Mini 3.8B (Q4)
(c) GPU server
Llama 3.3 70B (FP8)
Kimi Linear 48B-A3B (bf16)
Phi-4 Mini 3.8B (bf16)
Qwen 2.5 7B (AWQ)
Llama 3.2 1B (bf16)
Mean end-to-end latency (s)